

Introduction to Machine Learning

Homework Solution - 1

Instructor: Prof Chinmay Hegde

TA: Devansh Bisla, Maryam Majzoubi

Problem 1

Let $\{x_1, x_2, \dots, x_n\}$ be a set of points in d -dimensional space. Suppose we wish to produce a single point estimate $\mu \in \mathbb{R}^d$ that minimizes the squared-error:

$$\|x_1 - \mu\|_2^2 + \|x_2 - \mu\|_2^2 + \dots + \|x_n - \mu\|_2^2$$

Find a closed form expression for μ and prove that your answer is correct.

(Solution) We want to minimize the following error:

$$Error = \sum_{i=1}^n (x_i - \mu)^2$$

so we take the derivative w.r.t μ and set it to zero:

$$Error = -2 * \sum_{i=1}^n (x_i - \mu) = 0$$

This will result in following closed form solution, which is the mean of x_i 's:

$$\mu^* = \frac{\sum_{i=1}^n x_i}{n}$$

Problem 2

Not all norms behave the same; for instance, the ℓ_1 -norm of a vector can be dramatically different from the ℓ_2 -norm, especially in high dimensions. Prove the following norm inequalities for d -dimensional vectors, starting from the definitions provided in class and lecture notes. (Use any algebraic technique/result you like, as long as you cite it.)

- (a) $\|x\|_2 \leq \|x\|_1 \leq \sqrt{d}\|x\|_2$
- (b) $\|x\|_\infty \leq \|x\|_2 \leq \sqrt{d}\|x\|_\infty$
- (c) $\|x\|_\infty \leq \|x\|_1 \leq d\|x\|_\infty$

(Solution)

$$\|x\|_1 = \sum_{i=1}^d |x_i|, \quad \|x\|_2 = \sqrt{\sum_{i=1}^d |x_i|^2}, \quad \|x\|_\infty = \max\{|x_i| : i = 1, \dots, d\}$$

(a) For left hand side, we know that,

$$\sum_{i=1}^d |x_i|^2 \leq \left(\sum_{i=1}^d |x_i| \right)^2 \implies \|x\|_2^2 \leq \|x\|_1^2 \implies \|x\|_2 \leq \|x\|_1$$

For right hand side, according to Cauchy Swartz inequality, for any two d -dimensional vectors u and v ,

$$\left| \sum_{i=1}^d u_i v_i \right|^2 \leq \sum_{j=1}^d |u_j|^2 \sum_{k=1}^d |v_k|^2$$

Now, let $u_i = |x_i|$ and $v_i = 1 \forall i$ therefore,

$$\sum_{i=1}^d |x_i| \cdot 1 \leq \sqrt{\sum_{i=1}^d x_i^2} \sqrt{\sum_{i=1}^d 1^2} \implies \|x\|_1 \leq \sqrt{d}\|x\|_2$$

Hence, $\|x\|_2 \leq \|x\|_1 \leq \sqrt{d}\|x\|_2$

(b) The left hand side is evident from the definition itself, as the max element of a vector will be less than sum of squares of all the elements i.e

$$\left(\max\{|x_i| : i = 1, \dots, d\} \right)^2 \leq \sum_{i=1}^d |x_i|^2 \implies \|x\|_\infty \leq \|x\|_2$$

For right hand side note that, sum of the squares of the elements of a vector will always at best be equal to summing the largest element d times i.e

$$\sum_{i=1}^d |x_i|^2 \leq \sum_{i=1}^d \left(\max\{|x_i| : i = 1, \dots, d\} \right)^2 \implies \|x\|_2^2 \leq d\|x\|_\infty^2 \implies \|x\|_2 \leq \sqrt{d}\|x\|_\infty$$

(c) We can obtain the inequality from (a) and (b)

Problem 3

When we think of a Gaussian distribution (a bell-curve) in 1, 2, or 3 dimensions, the picture that comes to mind is a "blob" with a lot of mass near the origin and exponential decay away from the origin. However, the picture is very different in higher dimensions (and illustrates the counter-intuitive nature of high-dimensional data analysis). In short, we will show that *Gaussian distributions are like soap bubbles*: most of the mass is concentrated near a shell of a given radius, and is empty everywhere else.

- (a) Fix $d = 3$ and generate 10,000 random samples from the standard multi-variate Gaussian distribution defined in \mathbb{R}^d .
- (b) Compute and plot the histogram of Euclidean norms of your samples. Also calculate the average and standard deviation of the norms.
- (c) Increase d on a coarsely spaced log scale all the way up to $d = 1000$ (say $d = 50, 100, 200, 500, 1000$), and repeat parts (a) and (b). Plot the variation of the average and the standard deviation of Euclidean norm of the samples with increasing d .
- (d) What can you conclude from your plot from part (c)?
- (e) ****Bonus, not for grade.**** Mathematically justify your conclusion using a formal proof. You are free to use any familiar laws of probability, algebra, or geometry.

(Solution)

(e)** The expected value of squared Euclidean norm of a d dimensional Gaussian is as follows:

$$E[\|X\|_2^2] = E[X_1^2 + \dots + X_d^2] = E[X_1^2] + \dots + E[X_d^2] = \text{Var}(X_1) + \dots + \text{Var}(X_d) = d$$

Therefore, the expected value of Euclidean norm will converge to \sqrt{d} .

Thanks to Dimitrios Chaikalis for the following solution for (a - d):

HW1_Q3

February 4, 2020

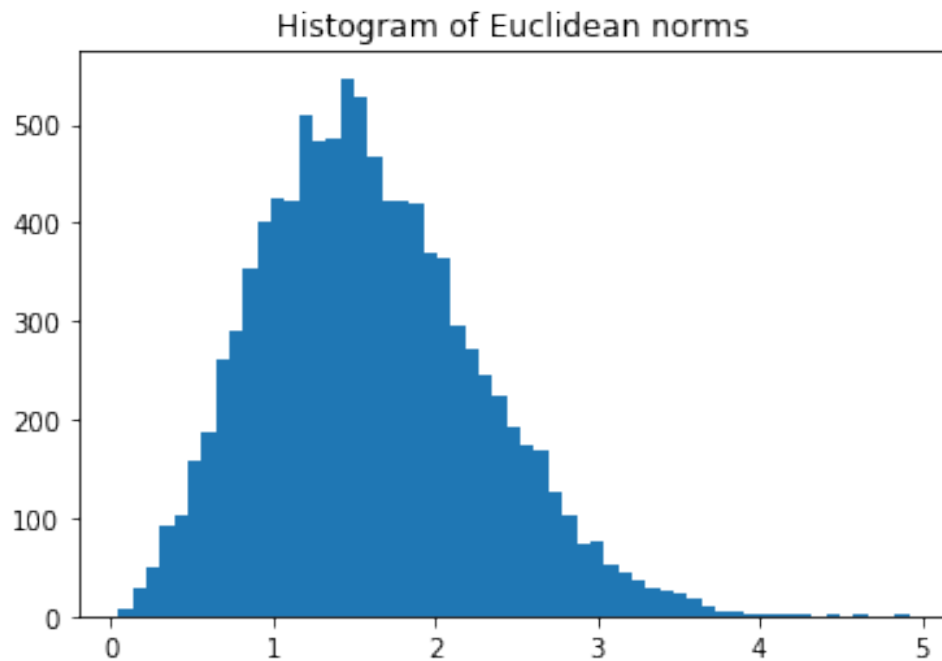
```
[0]: import numpy as np
import matplotlib.pyplot as plt

d = 3
means = np.zeros(d)
cov = np.eye(d)
x = np.random.multivariate_normal(means, cov, 10000)
```

Samples have been generated.

```
[0]: y = np.zeros(10000)
for i in range(10000):
    y[i] = np.linalg.norm(x[:,i], 2)
_ = plt.hist(y, bins='auto')
plt.title('Histogram of Euclidean norms')
```

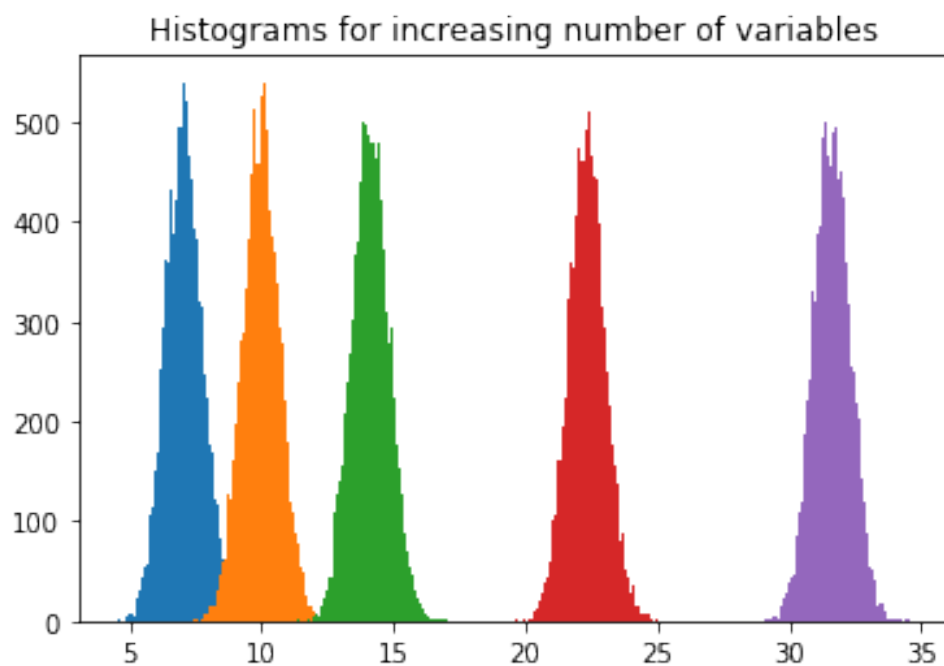
```
[0]: Text(0.5, 1.0, 'Histogram of Euclidean norms')
```



```
[0]: m = np.mean(y)
std = np.std(y)
print(m,std)
```

1.5983079784218295 0.6739420978301629

```
[0]: D = [50,100,200,500,1000]
M = np.zeros(5)
Std = np.zeros(5)
for i in range(5):
    means = np.zeros(D[i])
    cov = np.eye(D[i])
    x = np.random.multivariate_normal(means,cov,10000)
    y = np.zeros(10000)
    for j in range(10000):
        y[j] = np.linalg.norm(x[:,j],2)
    _ = plt.hist(y,bins='auto')
    M[i] = np.mean(y)
    Std[i] = np.std(y)
plt.title('Histograms for increasing number of variables')
```



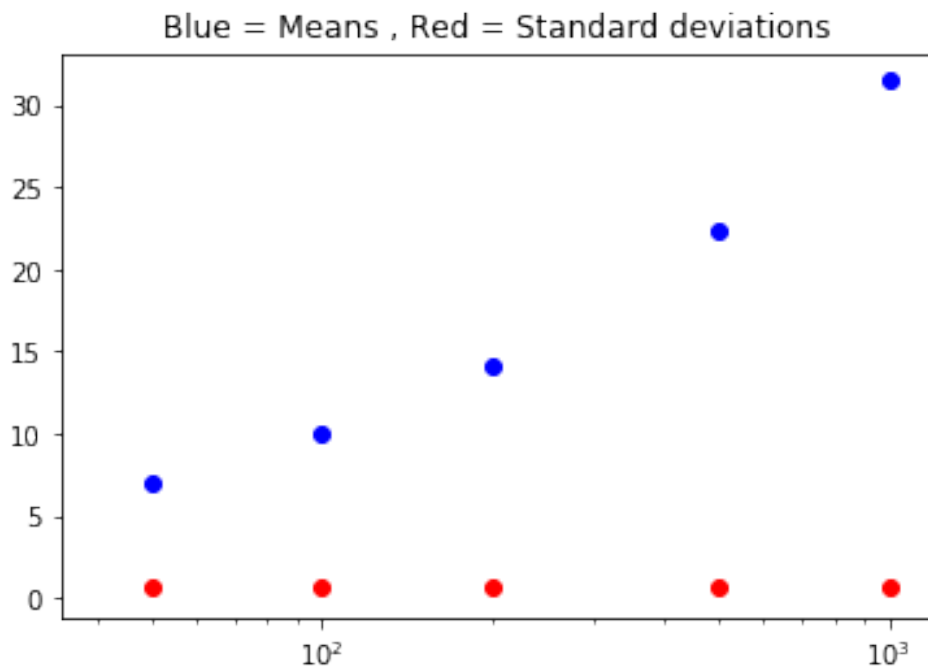
```
[0]: print(M)
```

```
[ 7.05201213  9.97963425 14.12533543 22.34299976 31.6195839 ]
```

```
[0]: print(Std)
```

```
[0.70726022 0.70237244 0.71003798 0.70606778 0.71128948]
```

```
[0]: _ = plt.scatter(D,M,color='blue')
_ = plt.scatter(D,Std,color='red')
plt.xscale('log')
plt.title('Blue = Means , Red = Standard deviations')
plt.show()
```



We can see that, apparently, as we increase the number of variables in the distributions, the norms of the resulting vectors increase in magnitude as well, apparently converging to the \sqrt{d} . Hence the mean value of the norms increases. Their deviation however remains unaffected.

Problem 4

The goal of this problem is to implement a very simple text retrieval system. Given (as input) a database of documents as well as a query document (all provided in an attached .zip file), write a program, in a language of your choice, to find the document in the database that is the best match to the query. Specifically:

- (a) Write a small parser to read each document and convert it into a vector of words.
- (b) Compute tf-idf values for each word in every document as well as the query.
- (c) Compute the cosine similarity between tf-idf vectors of each document and the query.
- (d) Report the document with the maximum similarity value.

(Solution) Thanks to Dimitrios Chaikalis for the following solution for (a - d):

HW1_Q4

February 8, 2020

```
[0]: import numpy as np
import math

d1 = open('d1.txt', 'r')
d2 = open('d2.txt', 'r')
d3 = open('d3.txt', 'r')
d4 = open('d4.txt', 'r')
d5 = open('d5.txt', 'r')
dq = open('d_query.txt', 'r')

def get_words(fileid):
    words = []
    temp = []
    while(True):
        r = fileid.read(1)
        if(r.isalpha()):
            temp.append(r)
        else:
            word = ''.join(temp)
            if(word!=""):
                in_lowercase = word.casefold()
                words.append(in_lowercase)
                temp = []
            if(r==' '):
                break
    return words

D1 = get_words(d1)
D2 = get_words(d2)
D3 = get_words(d3)
D4 = get_words(d4)
D5 = get_words(d5)
DQ = get_words(dq)
```

In the above cell, we read the text files, ignore all non-words (e.g. numbers and punctuation) then turn all words into lowercase to avoid ambiguities, and make a vector for each text file.


```
[0]: Dictionary = []
D = [D1 , D2 , D3 , D4 , D5 , DQ]
for i in range(6):
    for j in range(len(D[i])):
        temp = D[i][j]
        index = Dictionary.count(temp)
        if(index==0):
            Dictionary.append(temp)
```

Dictionary, meaning a vector with all the words appearing in the documents, is computed. Crucial for finding the ‘idf’ value and also dictates our vector length.

```
[0]: idf_map = {}
for i in range(len(Dictionary)):
    temp = Dictionary[i]
    count = 0
    for j in range(6):
        if(temp in D[j]):
            count = count + 1
    idf_map[temp] = math.log(6/count)
```

‘idf_map’ is a python dictionary structure, assigning every word its ‘idf’ value.

```
[0]: def get_tf_map(word_vector):
    tf_map = {}
    for i in range(len(word_vector)):
        temp = word_vector[i]
        if(temp in tf_map):
            tf_map[temp] = (tf_map[temp] + 1)
        else:
            tf_map[temp] = 1
    return tf_map

tf_1 = get_tf_map(D1)
tf_2 = get_tf_map(D2)
tf_3 = get_tf_map(D3)
tf_4 = get_tf_map(D4)
tf_5 = get_tf_map(D5)
tf_q = get_tf_map(DQ)
```

‘tf_i’ is a map assigning every word of document ‘i’ its ‘tf’ value.

```
[0]: def get_complete_vector(idf,tf):
    V = []
    for word in idf:
        if(word in tf):
            val = idf[word]*tf[word]
        else:
```

```

        val = 0
        V.append(val)
    return V

V1 = get_complete_vector(idf_map,tf_1)
V2 = get_complete_vector(idf_map,tf_2)
V3 = get_complete_vector(idf_map,tf_3)
V4 = get_complete_vector(idf_map,tf_4)
V5 = get_complete_vector(idf_map,tf_5)
VQ = get_complete_vector(idf_map,tf_q)

```

'Vi' is the final vector for document 'i', containing the 'tf*idf' value of each word in the appropriate cell.

```

[0]: def get_cosine(vec,q_vec):
        dot = np.dot(vec,q_vec)
        nv = np.linalg.norm(vec,2)
        nq = np.linalg.norm(q_vec,2)
        cosine = dot/(nv*nq)
        return cosine

c1 = get_cosine(V1,VQ)
c2 = get_cosine(V2,VQ)
c3 = get_cosine(V3,VQ)
c4 = get_cosine(V4,VQ)
c5 = get_cosine(V5,VQ)

```

Finally, we calculate the cosine similarity, and find which document gets the greatest value.

```

[0]: C = [c1,c2,c3,c4,c5]
num = np.argmax(C) + 1
print('The best match, is document d%d.txt'%num)

```

The best match, is document d4.txt