


Intro to ML Homework5

Haotian Yi N18800809

April 14, 2020

1 Question 1

Computation graph:



Forward pass:

$$z = w_1 * x$$

$$u = \sigma(z)$$

$$\phi = w_2 * u$$

$$L = (\phi - y)^2$$

Backward pass:

$$\partial_L L = 1$$

$$\partial_\phi L = 2(\phi - y)$$

$$\partial_{w_2} \phi = u$$

$$\partial_u \phi = w_2$$

$$\partial_z u = \sigma'(z) = \sigma(z)(1 - \sigma(z)) = u(1 - u)$$

$$\partial_{w_1} z = x$$

Thus:

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial \phi} \frac{\partial \phi}{\partial u} \frac{\partial u}{\partial z} \frac{\partial z}{\partial w_1} = 2(\phi - y) * w_2 * u(1 - u) * x$$

$$\frac{\partial L}{\partial w_2} = \frac{\partial L}{\partial \phi} \frac{\partial \phi}{\partial w_2} = 2(\phi - y) * u$$

It can be seen that forward pass uses one sigmoid, two multiplication, one subtraction and a square, backward pass calculate multiple derivatives, it looks like backward pass is more complex by almost twice. However when we calculate derivative respectively on w_1 or w_2 , we only need four multiplication on what we got by derivatives, but there would be several duplicate calculations (as shown below) in forward pass which contains several multiplication, several complex sigmoid, these will lower its efficiency.

$$\frac{\partial L}{\partial w_1} = 2(w_2\sigma(z) - y) * w_2 * \sigma(z)(1 - \sigma(z)) * x$$

$$\frac{\partial L}{\partial w_2} = 2(w_2\sigma(z) - y) * \sigma(z)$$

Roughly I think backward pass used for calculating gradient will save half calculation than forward method.

2 Question 2

(a) Z's shape: (46,62,20), U's shape: (46,62,20)

(b) number of input channel: 10, number of output channel: 20.

(c)

Add-operation: $46*62*9*10*2 = 513360$

Multiplication: $46*62*9*10 = 256680$

$(46*62*20)*(3*3*10) = 5133600$
add and multiply-operations

(d) Number of trainable parameters is $:3*3*10*20 = 1800$

$(3*3*10*20)+20 = 1820$

3 Question 3

(a) Regularized loss:

$$L = L(w) + \lambda \|w\|$$

(b) Update Rule:

$$w_{t+1} = w_t - \alpha \frac{\partial L}{\partial w} = w_t - \alpha * \left(\frac{\partial L(w)}{\partial w} + 2\lambda w_t \right) = (1 - 2\alpha\lambda)w_t - \alpha * \frac{\partial L(w)}{\partial w}$$

(c) As we can see from (b), before each descent update, w_t is multiplied by $(1 - 2\alpha\lambda)$, which is a smaller factor than 1, so that it make weights shrunk.

(d) Increasing λ will lift up strength of regularization, it makes w shrink more intensively. Learning rate should be chosen as a small number in the range of 0 to 1, and for the update above, I think learning rate should be chosen to make $2\alpha\lambda$ much smaller than 1 to make the update stable.

4 Question 4

upload data and unzip



```
!unzip data.zip
```

import model and normalize data

```
[18] import torchvision
import torch
from torchvision import datasets, models, transforms
transforms = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize([0.485, 0.456, 0.406],
        [0.229, 0.224, 0.225])
])
train_set = datasets.ImageFolder("data/train", transforms)
val_set = datasets.ImageFolder("data/val", transforms)
model = models.resnet34(pretrained=True)

trainloader = torch.utils.data.DataLoader(train_set, batch_size=4,
                                           shuffle=True, num_workers=2)
valloader = torch.utils.data.DataLoader(val_set, batch_size=1,
                                         shuffle=True, num_workers=2)
classes = ('cat', 'dog')
```

take a look at train set

```
import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2.5 + 0.25      # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = dataiter.next()
print('labels:', labels)
# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join('%5s' % classes[labels[j]] for j in range(4)))
```

➡ Clipping input data to the valid range for imshow with RGB data ([0..1
labels: tensor([1, 1, 1, 0])



modify last layer for 2 classes output

```
[20] import torch.nn as nn
      model.fc = nn.Linear(model.fc.in_features, 2)
```

define loss function and optimizer

```
[21] import torch.optim as optim

      criterion = nn.CrossEntropyLoss()
      optimizer = optim.SGD(model.parameters(), lr=0.001)
```

train the network with validation

```
[22] avgloss = []
      total = 0
      correct = 0
      vloss = []
      vcorrect = 0
      vtotal = 0
      for epoch in range(20): # loop over the dataset multiple times
          # one epoch means run over whole data once
          #####Train#####
          running_loss = 0.0
          for i, data in enumerate(trainloader,0):
              #print(i)
              # get the inputs; data is a list of [inputs, labels]
              inputs, labels = data
              #print(labels)
              # zero the parameter gradients
              optimizer.zero_grad()

              # forward + backward + optimize
              outputs = model(inputs)

              _, predicted = torch.max(outputs, 1)
              #print(predicted)
              for j in range(4):
                  total = total + 1
                  if labels[j] == predicted[j]:
                      correct = correct + 1;
```

```

        #print(correct,'of',total, 'are correctly predicted')

        loss = criterion(outputs, labels)
        #print('every loss:',loss)
        loss.backward()
        optimizer.step()

        # print statistics
        running_loss += loss.item()
        # print every 15 mini-batches (15*4 = 16 = #pictures in one epoch)
        print(correct,'of',total, 'are correctly predicted')
        print('For epoch %d, every %5d mini-batches, average loss: %.3f' % (epoch + 1, 15, running_loss / 15))
        avgloss.append(running_loss / 15)
        running_loss = 0.0
#####Validation#####
    with torch.no_grad():
        dataiter = iter(valloader)
        vimages, vlabels = dataiter.next()
        vimages, vlabels = data
        voutputs = model(vimages)
        VLOSS = criterion(voutputs, vlabels)
        vloss.append(VLOSS)
        _, vpredicted = torch.max(voutputs, 1)
        vttotal += vlabels.shape[0]
        vcorrect += (vpredicted == vlabels).sum().item()
        print('Validation loss: ',VLOSS.item(),'; ',vcorrect,'of',vttotal,'are correctly predicted')

```

It will print the process like below. There is just a part of the process, it can be fully viewed in attached python file.

```

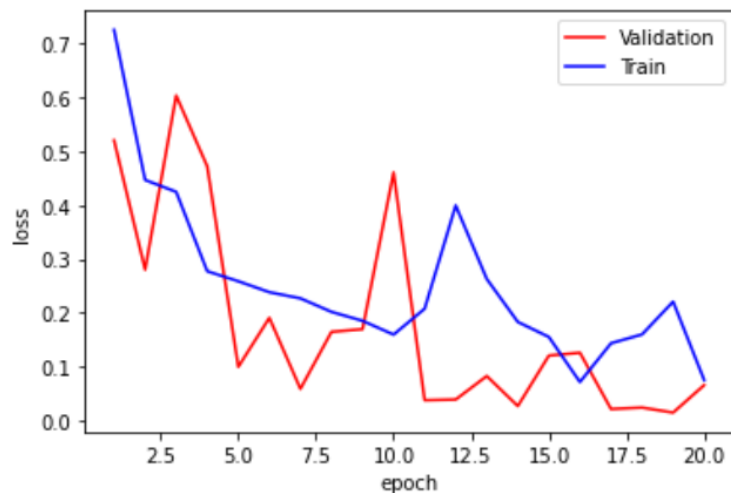
For epoch 14, every    15 mini-batches, average loss: 0.183
Validation loss:  0.026706388220191002 ;  53 of 56 are correctly predicted
803 of 900 are correctly predicted
For epoch 15, every    15 mini-batches, average loss: 0.155
Validation loss:  0.1204538494348526 ;  57 of 60 are correctly predicted
863 of 960 are correctly predicted
For epoch 16, every    15 mini-batches, average loss: 0.071
Validation loss:  0.12573614716529846 ;  61 of 64 are correctly predicted
921 of 1020 are correctly predicted
For epoch 17, every    15 mini-batches, average loss: 0.143
Validation loss:  0.021214816719293594 ;  65 of 68 are correctly predicted
978 of 1080 are correctly predicted
For epoch 18, every    15 mini-batches, average loss: 0.160
Validation loss:  0.02385798469185829 ;  69 of 72 are correctly predicted
1031 of 1140 are correctly predicted
For epoch 19, every    15 mini-batches, average loss: 0.221
Validation loss:  0.014515489339828491 ;  73 of 76 are correctly predicted
1091 of 1200 are correctly predicted
For epoch 20, every    15 mini-batches, average loss: 0.075
Validation loss:  0.0656055212020874 ;  77 of 80 are correctly predicted

```

Plot loss curves for train and validation

```
print('Finished Training')
print('Accuracy of the validation images: %d %%' % (100 * vcorrect / vtotal))
print('Accuracy of the the test images: %d %%' % (100 * correct / total))
plt.figure()
vepoch = range(1,21)
plt.plot(vepoch,vloss,'r',label='Validation')
plt.legend()
epoch = range(1,21)
plt.plot(epoch,avgloss,'b',label='Train')
plt.legend()
plt.xlabel('epoch')
plt.ylabel('loss')
```

```
Finished Training
Accuracy of the validation images: 96 %
Accuracy of the the test images: 90 %
Text(0, 0.5, 'loss')
```



Loss curves and accuracy are shown above. Accuracy of validation is 96%, Accuracy of training is 90% .

I also tried a test use all pictures in validation directory for curiosity and display accuracy for each class, it can be checked in python document attached.