

# Intro to ML Homework6

Haotian Yi N18800809

April 29, 2020

## 1 Question 1

```
[ ] import numpy as np
X = np.array([[3,2,1],[2,4,5],[1,2,3],[0,2,5]])
print(X)
```

```
↳ [[3 2 1]
    [2 4 5]
    [1 2 3]
    [0 2 5]]
```

a. Find the sample mean

```
[ ] X_mean = np.mean(X, axis = 0)
print(X_mean)
```

```
↳ [1.5 2.5 3.5]
```

b. Zero-center the samples, and find the eigenvalues and eigenvectors of the data covariance matrix Q.

(1) Zero-center the samples

```
[ ] X_original = X.copy()
X = X - X_mean
print(X)
```

```
↳ [[ 1.5 -0.5 -2.5]
    [ 0.5  1.5  1.5]
    [-0.5 -0.5 -0.5]
    [-1.5 -0.5  1.5]]
```

(2) Find the eigenvalues and eigenvectors of the data covariance matrix Q.

```
▶ Q = 1/4 * np.matmul(X.T,X) # covariance matrix Q
value,evector = np.linalg.eig(Q)
print(value) # eigenvalues #print(evector) # eigenvectors(a row is a vector)
evector1 = np.array(evector[:,0])
evector2 = np.array(evector[:,1])
evector3 = np.array(evector[:,2])
evec = np.array([evector1,evector2,evector3])
print(evec)
```

```
↳ [3.56166464 1.1733803 0.01495506]
[[-0.45056922 0.19247228 0.87174641]
 [-0.66677184 -0.72187235 -0.18524476]
 [-0.59363515 0.66472154 -0.45358856]]
```

First list is of eigenvalues, each row of the matrix above is an eigenvector.

c. Find the PCA coefficients corresponding to each of the samples in X.

```
[ ] scores = np.matmul(X,evec.T)
    print(scores)

[ ] [[-2.95145599 -0.17610969 -0.0888421 ]
     [ 1.37104342 -1.69406159  0.0198819 ]
     [-0.30682473  0.78694448  0.19125108]
     [ 1.8872373   1.0832268  -0.12229089]]
```

PCA coefficients(scores) produced by eigenvector for each samples are shown above, values in each row are three scores of each sample.

d. Reconstruct the original samples from the top two principal components, and report the reconstruction error for each of the samples.

PCA reconstruction= PC scores \* Eigenvectors<sup>T</sup> + Mean

```
[ ] np.set_printoptions(precision=4)
    eigen12 = np.delete(evec, 2, 0)
    #print(eigen12)
    PC12 = np.delete(scores, 2, 1)
    #print(PC12)
    recons = np.matmul(PC12,eigen12)+ X_mean
    print(recons)

[ ] [[ 2.9473  2.0591  0.9597]
     [ 2.0118  3.9868  5.009 ]
     [ 1.1135  1.8729  3.0867]
     [-0.0726  2.0813  4.9445]]
```

Reconstruction error

```
[ ] #print(recons-X_original)
    RE = np.linalg.norm((recons-X_original),axis=1)
    print(RE)

[ ] [0.0888 0.0199 0.1913 0.1223]
```

Reconstruction errors for each sample are shown above in the output list.

## 2 Question 2

(a)

Loss function of k-means is:

$$F(\{S_1, \dots, S_k\}, \{\mu_1, \dots, \mu_k\}) = \sum_{j=1}^k \sum_{x_i \in S_j} \|x_i - \mu_j\|^2$$

We want it in the form:

$$F(\eta, \mu) = \sum_{i=1}^n \sum_{j=1}^k \eta_{ij} \|x_i - \mu_j\|^2$$

So we want  $\sum_{i=1}^n \eta_{ij}$  serves the function of  $\sum_{x_i \in S_j}$ ,  $\sum_{i=1}^n \sum_{j=1}^k \eta_{ij}$  bonds  $x_i$  and corresponding  $\mu_j$ , so  $\eta$  is a binary matrix of size of (n,k), for its element, for instance, if a  $x_i$ 's corresponding center is  $\mu_j$ , then  $\eta_{ij}$  is 1, rest  $\eta_i$  in the same row of the matrix is 0.

(b)

Lloyd's algorithm can only decrease the value of  $F$ , because at each iteration:

(1) it re-arrange points to center which is closest to them, and then compute new centers as mean of each cluster,

$$(2) F(\{S_1, \dots, S_k\}, \{\mu_1, \dots, \mu_k\}) = \sum_{j=1}^k \sum_{x_i \in S_j} \|x_i - \mu_j\|^2 = \sum_{j=1}^k |S_j| * \text{Variance} S_j,$$

(3) set centers to mean points  $\mu_j$  in  $S_j$ . This is equivalent to minimizing the pairwise squared deviations of points in the same cluster in the function above, thus  $F$  is decreased.

(c)

The algorithm will terminate in no more than T iterations, where T is some finite number, because at each iteration of Lloyd's algorithm, it will reduce the value of  $F$ , and it can not iterate infinitely which means  $F$  can be reduced infinitely.

For upper bound of T, considering there are n points and k centers, there could be  $k^n$  combination of points and their corresponding centers, assume the worst case: only one point changes its corresponding center and algorithm has to try  $k^n$  combination and find the best combination at last. So the upper bound of T is  $O(k^n)$ .

### 3 Question 3

Using the Senate Votes dataset demo'ed in Lecture 11, perform k-means clustering with  $k = 2$  and show that you can learn (most of) the Senators' parties in a completely unsupervised manner. Which Senators did your algorithm make a mistake on, and why?

(1) Import dataset

```
[ ] import numpy as np
import matplotlib
import matplotlib.pyplot as plt
```

```
[ ] import pandas as pd
```

```
url_votes = 'https://raw.githubusercontent.com/exemplary-citizen/PCA-and-Senate-Voting-Data/master/senator_pca_problem/senator_data_p
senator_df = pd.read_csv(url_votes + 'data_matrix.csv', error_bad_lines=False)
af = pd.read_csv(url_votes + 'politician_labels.txt', header=None)
af["affiliations"] = af[0].str.split().str[-1]
```

```
[ ] senator_df.head()
```



	Unnamed: 0	bill_type	bill_name	bill_ID	missing_votes	Arlen Specter (PA)	Barack H. Obama (IL)	Barbara Mikulski (MD)	Barbara Boxer (CA)	Bill Nelson (FL)	Blanche Lincoln (AR)
0	0	Appropriations	Transit Security Amendment_3866		0.0	1.0	1.0	1.0	1.0	1.0	
1	1	Budget_Spending_and_Taxes_2007	Budget Resolution...		0.0	1.0	-1.0	-1.0	-1.0	-1.0	
2	2	"Budget, Spending and Taxes	Debt Limit Increases...		0.0	1.0	-1.0	-1.0	-1.0	-1.0	
3	3	"Budget, Spending and Taxes	Education Funding ...		0.0	-1.0	1.0	1.0	1.0	1.0	
4	4	Budget, Spending and Taxes	Reinstate Pay-As-You...		0.0	-1.0	1.0	1.0	1.0	1.0	

5 rows × 103 columns

```
[ ] af.head()
```



	0	affiliations
0	Specter Red	Red
1	Obama Blue	Blue
2	Mikulski Blue	Blue
3	Boxer Blue	Blue
4	Nelson Blue	Blue

```
[ ] X = np.array(senator_df.values[:, 3:].T, dtype='float64') #transpose to get senators as rows
```

## (2) Check shape of matrix, affiliations

Check shape of data sets

```
[ ] print(senator_df.shape)
    print(af.shape)
    #print(af)
    #print(af["affiliations"])
    print(X.shape)
    print(X)
```

```
↳ (542, 103)
   (100, 2)
   (100, 542)
   [[ 1.  1.  1. ...  1.  1.  1.]
    [ 1. -1. -1. ...  0.  0.  0.]
    [ 1. -1. -1. ...  1. -1. -1.]
    ...
    [-1.  1.  1. ...  1.  1.  1.]
    [-1.  1.  1. ...  0.  0.  0.]
    [-1.  1.  1. ...  1.  1.  1.]]
```

Check part of affiliations


```
[ ] affiliations = af["affiliations"]
    print(affiliations[30:40])
```

```
↳ 30      Blue
   31      Blue
   32       Red
   33       Red
   34     Yellow
   35      Blue
   36       Red
   37       Red
   38       Red
   39       Red
   Name: affiliations, dtype: object
```

## (3) Visualize affiliations

Visualize affiliations

```
[50] affiliations = af["affiliations"]
     senator_num = range(100)
     plt.figure(figsize=(18,5))
     plt.scatter(senator_num, np.zeros_like(senator_num), c=affiliations)
     plt.title('Parties of senators')
     plt.show()
```



#### (4) Perform K-means

Perform k-means clustering with  $k = 2$  and show that you can learn (most of) the Senators' parties in a completely unsupervised manner.

```
[51] from sklearn.cluster import KMeans
      kmeans = KMeans(2)
      labels = kmeans.fit(X).predict(X)
      print(labels)

[0 1 1 1 1 1 1 0 1 0 1 0 1 0 0 0 0 1 1 0 1 1 0 1 1 0 1 1 0 0 0 1 1 1 0 0 0 1 1 0
 0 0 0 1 0 0 1 1 0 0 0 0 0 0 1 1 0 0 1 1 0 0 1 0 0 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0
 0 1 1 1 0 0 0 1 0 0 1 0 1 1 1 0 0 0 0 0 0 1 1 1 0 0 0]
```


```
[54] af_kmeans = []
      mislabel = []
      yellowbabels = []
      for i in range(100):
          if labels[i] == 1:
              af_kmeans.append('Blue')
          elif labels[i] == 0:
              af_kmeans.append('Red')
          if af_kmeans[i] != affiliations[i]:
              mislabel.append(i)
          if affiliations[i] == 'Yellow':
              yellowbabels.append(i)
      print("Senators who are mis labelled: ",mislabel)
      print("Among them, these(Yellow) are neither of each label: ",yellowbabels)
      for i in mislabel:
          print(af.values[i][0])
```

```
Senators who are mis labelled: [21, 34, 58, 62]
Among them, these(Yellow) are neither of each label: [34, 62]
Nelson Blue
Jeffords Yellow
Chafee Red
Dayton Yellow
```

Senators who are mislabeled are shown above in the output.

#### (5) Visualize the result of k-means

```
[58] plt.figure(figsize=(18,5))
      plt.scatter(senator_num, np.zeros_like(senator_num), c=af_kmeans, s=40, cmap='viridis');
      plt.plot(mislabel,np.zeros_like(mislabel),'cs',label='mislabeled')
      plt.plot(yellowbabels,np.zeros_like(yellowbabels),'ys',label='mislabeled from yellow')
      plt.legend()
      plt.show()
```



Square points are mislabeled senators whose original label is red/blue, yellow points are senators who are originally of yellow label. Which Senators did your algorithm make a mistake on, and why? **21,34,58,62.**

There are two senator originally labeled with 'Yellow', these will inevitably be wrong labeled due to only two clusters. And there are many potential drawbacks of k-means leading to rest two senators mislabeled: maybe these two are on the boundary of two clusters and closer to wrong cluster thus they are easily mislabeled, maybe distribution is not spherical.

## 4 Question 4

There is problem loading The Places Rated Almanac ratings from txt file provide so I find same csv file online.

a. Load the data and construct a table with 9 columns containing the numerical ratings.

import data

```
[2] import pandas as pd

Ratings = pd.read_csv('places.csv', error_bad_lines=False)
```

```
[ ] Ratings.head()
```

	Place	Climate_and_Terrain	Housing	Health_Care_and_Environment	Crime	Transportation	Education	The_Arts	Recreation	Economics
0	Abilene_TX	521	6200	237	923	4031	2757	996	1405	7633
1	Akron_OH	575	8138	1656	886	4883	2438	5564	2632	4350
2	Albany_GA	468	7339	618	970	2531	2560	237	859	5250
3	Albany-Schenectady-Troy_NY	476	7908	1431	610	6883	3399	4655	1617	5864
4	Albuquerque_NM	659	8393	1853	1483	6558	3026	4496	2612	5727

Code

Text

Extract places and ratings

```
[3] import numpy as np

place = np.array(Ratings.values[:,0])
print(place[0])
print(place.shape)

ratings = np.array(Ratings.values[:,1:], dtype='float64')
print(ratings[0,:])
print(ratings.shape)
```

```
Abilene_TX
(329,)
[ 521. 6200. 237. 923. 4031. 2757. 996. 1405. 7633.]
(329, 9)
```

b. Replace each value in the matrix by its base-10 logarithm. (This pre-processing is done for convenience since the numerical range of the ratings is large.)

```
[4] ratings_original = ratings
     ratings = np.log10(ratings)
     print(ratings[0,:])
     print(ratings.shape)
```

```
[2.71683772 3.79239169 2.37474835 2.9652017  3.6054128  3.44043677
 2.99825934 3.14767632 3.88269526]
(329, 9)
```

c. Perform PCA on the data. Remember to center the data points first by computing the mean data vector and subtracting it from every point. With the centered data matrix, do an SVD and compute the principal components.

```
[5] from sklearn.decomposition import PCA

     X = ratings
     X_mean = np.mean(X, axis = 0)
     X_zc = X - X_mean
     pca = PCA(n_components=2)
     pca.fit(X_zc)
     e_vectors = pca.components_ # get eigenvectors
```

d. Write down the first two principal components v1 and v2. Provide a qualitative interpretation of the components. Which among the nine factors do they appear to correlate the most with?

```
[6] print('eigenvectors:')
     print(e_vectors)
```

```
eigenvectors:
[[ 0.03507288  0.09335159  0.40776448  0.10044536  0.15009714  0.03215319
   0.87434057  0.15899622  0.01949418]
 [ 0.0088782  0.00923057 -0.85853187  0.22042372  0.05920111 -0.06058858
   0.30380632  0.33399255  0.0561011 ]]
```

Principal components(eigenvectors) are shown above, these vectors point the direction data varies the most.


For first component, 3rd and 7th factors contribute more for it, we can say they are correlated, also we can see 7th factor is much larger than rest, it is dominant, in fact, we could state that based on the correlation of 0.87434057 that this principal component is primarily a measure of the Arts. 3rd factor is Health Care and Environment, 7th factor is The Arts.

For second component, 3rd factor contributes the most for it. It seems 3rd factor dominate the vector and have small correlation with other unimportant factors because -0.85853187 is much larger than others. 3rd factor is Health Care and Environment.



e. Project the data points onto the first two principal components. (That is, compute the highest 2 scores of each of the data points.) Plot the scores as a 2D scatter plot. Which cities correspond to outliers in this scatter plot?

```
import matplotlib.pyplot as plt
scores = np.matmul(X_zc, e_vectors.T)
s_mean = np.mean(scores, axis=0)
s_std = np.std(scores, axis=0)
outliers = []
for i in range(scores.shape[0]):
    if (scores[i][0] > s_mean[0]+3*s_std[0] or scores[i][0] < s_mean[0]-3*s_std[0]) or (scores[i][1] > s_mean[1]+3*s_std[1] or scores[i][1] < s_mean[1]-3*s_std[1]):
        outliers.append(i)
plt.figure(figsize=(6.5,6.5))
plt.scatter(scores[:,0], scores[:,1], c='b')
for i in outliers:
    plt.plot(scores[i][0], scores[i][1], 'ro')
```



For outliers, it is decided by  $3\sigma$  threshold from Gaussian Distribution. Red points above are outliers.

Print outliers:

```
for i in outliers:
    print(place[i])
```

Brownsville-Harlington\_TX  
Midland\_TX  
New\_York\_NY

Cites which are outliers are shown above.

f. Repeat Steps 2-5, but with a slightly different data matrix – instead of computing the base-10 logarithm, use the z-scores. (The z-score is calculated by computing the mean  $\mu$  and standard deviation  $\sigma$  for each feature, and normalizing each entry  $x$  by  $(x-\mu)/\sigma$ ). How do your answers change?

(1) compute z

```
[30] np.set_printoptions(precision=4,suppress=True,floatmode = 'fixed')
      r_mean = np.mean(ratings_original,axis=0)
      r_std = np.std(ratings_original,axis=0)
      z = ratings_original
      print(ratings_original.shape)
      for i in range(329):
          for j in range(9):
              z[i][j] = (z[i][j] - r_mean[j]) / r_std[j]
      print(z)
```

```
(329, 9)
[[-0.1470 -0.9013 -0.9473 ... -0.4649 -0.5466  1.9464]
 [ 0.3007 -0.0876  0.4696 ...  0.5206  0.9744 -1.0855]
 [-0.5864 -0.4231 -0.5669 ... -0.6286 -1.2235 -0.2543]
 ...
 [ 0.0105  0.0103 -0.4720 ... -0.4593 -1.2446 -0.5351]
 [ 0.2592 -0.5566 -0.0886 ... -0.0763 -0.6433 -1.5066]
 [ 0.5742 -0.1980 -0.9723 ... -0.6534 -1.1504 -0.7678]]
```

data points are already centered

(2) Write down the first two principal components v1 and v2. Provide a qualitative interpretation of the components.

```
[31] from sklearn.decomposition import PCA

      pca = PCA(n_components=2)
      pca.fit(z)
      e_vectors = pca.components_ # get eigenvectors
      print('eigenvectors:')
      print(e_vectors)
```

```
eigenvectors:
[[ 0.2064  0.3565  0.4602  0.2813  0.3512  0.2753  0.4631  0.3279  0.1354]
 [ 0.2178  0.2506 -0.2995  0.3553 -0.1796 -0.4834 -0.1948  0.3845  0.4713]]
```

Principal components(eigenvectors) are shown above, these vectors point the direction data varies the most.


For first component, factor values are small and close to each others, we can relatively see that 3rd and 7th factors contribute more for it, we can say they are correlated. 3rd factor is Health Care and Environment, 7th factor is The Arts.

For second component, factor values are small and close to each others, we can relatively see that 6th and 9th factors contribute more for it, we can say they are correlated. 6th factor is Education, 9th factor is Economics.

(3) Project the data points onto the first two principal components. (That is, compute the highest 2 scores of each of the data points.) Plot the scores as a 2D scatter plot. Which cities correspond to outliers in this scatter plot?

```
[35] import matplotlib.pyplot as plt
      scores = np.matmul(z,e_vectors.T)
      s_mean = np.mean(scores,axis=0)
      s_std = np.std(scores,axis=0)
      outliers = []
      for i in range(scores.shape[0]):
          if (scores[i][0] > s_mean[0]+3*s_std[0] or scores[i][0] < s_mean[0]-3*s_std[0]) or (scores[i][1] > s_mean[1]+3*s_std[1] or scores[i][1] < s_mean[1]-3*s_std[1]):
              outliers.append(i)
      print(outliers)
      plt.figure(figsize=(6.5,6.5))
      plt.scatter(scores[:,0],scores[:,1],c='b')
      for i in outliers:
          plt.plot(scores[i][0],scores[i][1], 'ro')
```

☞ [42, 64, 167, 178, 212, 269, 313]



For outliers, it is decided by  $3\sigma$  threshold from Gaussian Distribution. Red points above are outliers.

Print outliers

```
[39] for i in outliers:
      print(place[i])
```

☞ Boston\_MA  
Chicago\_IL  
Las\_Vegas\_NV  
Los\_Angeles\_Long\_Beach\_CA  
New\_York\_NY  
San\_Francisco\_CA  
Washington\_DC-MD-VA

Cites which are outliers are shown above.