

demo3

February 26, 2020

This module explains the basics of optimization in ML and includes a simple implementation of gradient descent.

```
[0]: %matplotlib inline
import numpy as np
import matplotlib
import matplotlib.pyplot as plt
np.random.seed(1337)

kwargs = {'linewidth' : 3.5}
font = {'weight' : 'normal', 'size' : 24}
matplotlib.rc('font', **font)

def error_plot(ys, yscale='log'):
    plt.figure(figsize=(5, 5))
    plt.xlabel('Step')
    plt.ylabel('Error')
    plt.yscale(yscale)
    plt.plot(range(len(ys)), ys, **kwargs)
```

Let's now implement GD in one line! This is probably not memory efficient but useful to visualize what is going on.

```
[0]: def gradient_descent(init, steps, grad):
    xs = [init]
    for step in steps:
        xs.append(xs[-1] - step * grad(xs[-1]))
    return xs
```

Toy example.

```
[0]: def quadratic(x):
    return 0.5*(x-7).dot(x-7)

def quadratic_gradient(x):
    return x-7
```

Test your GD optimizer!

```
[0]: x0 = np.random.normal(0,1,(1000))
_, xout = gradient_descent(x0,[1.0],quadratic_gradient)
```


OK! looks like this ran quickly. In fact, for simple parabolic functions like this, GD optimizes in 1 step.

```
[0]: xout.shape
```

```
[0]: (1000,)
```

However, if we choose a smaller learning rate, things take longer.

```
[0]: xs = gradient_descent(x0, [0.1]*500, quadratic_gradient)
error_plot([quadratic(x) for x in xs])
```



Alright! Things seem to be converging. Lets now do linear regression.


```
[0]: def MSE(X,y,w):
      return 0.5*np.linalg.norm(X.dot(w)-y)**2
```

```
def MSE_gradient(X,y,w):  
    return X.T.dot(X.dot(w)-y)
```

We will cook up some fake data where the optimal regressor is known, so that we can compare errors.


```
[0]: n,d = 100,10  
X = np.random.normal(0,1,(n,d))  
wtrue = np.random.normal(0,1,d)  
noise = np.random.normal(0,0.01,n)  
y = X.dot(wtrue) + noise  
  
objective = lambda w: MSE(X,y,w)  
gradient = lambda w: MSE_gradient(X,y,w)  
  
w0 = np.random.normal(0,1,d)  
ws = gradient_descent(w0,[0.001]*300,gradient)  
error_plot([objective(w) for w in ws])  
plt.plot(range(len(ws)),[MSE(X,y,wtrue)]*len(ws))
```

```
[0]: [<matplotlib.lines.Line2D at 0x7f1905e9b4a8>]
```



Neat! The final MSE we get is very close to the noise level, which is basically as good as we can hope to do.

```
[0]: error_plot([np.linalg.norm(w-wtrue)**2 for w in ws])
```



Play around with various parameters (step size, epochs, etc) to see what happens.