

demo10_b

April 7, 2020

```
[0]: !pip install torchtext==0.4
```

```
Requirement already satisfied: torchtext==0.4 in /usr/local/lib/python3.6/dist-packages (0.4.0)
Requirement already satisfied: requests in /usr/local/lib/python3.6/dist-packages (from torchtext==0.4) (2.21.0)
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from torchtext==0.4) (1.12.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from torchtext==0.4) (1.18.2)
Requirement already satisfied: torch in /usr/local/lib/python3.6/dist-packages (from torchtext==0.4) (1.4.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.6/dist-packages (from torchtext==0.4) (4.38.0)
Requirement already satisfied: idna<2.9,>=2.5 in /usr/local/lib/python3.6/dist-packages (from requests->torchtext==0.4) (2.8)
Requirement already satisfied: urllib3<1.25,>=1.21.1 in /usr/local/lib/python3.6/dist-packages (from requests->torchtext==0.4) (1.24.3)
Requirement already satisfied: chardet<3.1.0,>=3.0.2 in /usr/local/lib/python3.6/dist-packages (from requests->torchtext==0.4) (3.0.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.6/dist-packages (from requests->torchtext==0.4) (2019.11.28)
```

The AG_NEWS dataset consists of 4 types of news articles (World, sports, business, sci/tech). 120K train samples and 76K test samples.

```
[0]: import torch
import torchtext
from torchtext.datasets import text_classification
import os
if not os.path.isdir('./.data'):
    os.mkdir('./.data')
train_dataset, test_dataset = text_classification.DATASETS['AG_NEWS'](root='./.data', ngrams=2, vocab=None)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
BATCH_SIZE=16
```

120000lines [00:06, 18333.38lines/s]

```
120000lines [00:13, 8919.15lines/s]
7600lines [00:00, 9705.19lines/s]
```

The dataset is in a TorchText object. We can query properties using inbuilt routines.

```
[0]: vocab_size = len(train_dataset.get_vocab())
      num_class = len(train_dataset.get_labels())
      print(vocab_size, num_class)
```

```
1308844 4
```

```
[0]: ag_news_label = {0 : "World",
                      1 : "Sports",
                      2 : "Business",
                      3 : "Tech"}
```

We need to write a method to get batches of data from this object. We collect bunch of text sequences (of variable length). Since lengths are variable, we store a tensor of delimiters (‘offsets’) to denote beginning of each sequence.

```
[0]: def generate_batch(batch):
      labels = torch.tensor([entry[0] for entry in batch])
      text = [entry[1] for entry in batch]
      offsets = [0] + [len(entry) for entry in text]
      offsets = torch.tensor(offsets[:-1]).cumsum(dim=0)
      text = torch.cat(text)
      return text, offsets, labels
```

Let’s train a simple language model to predict the type of news article. We will use an “embedding-bag” layer which converts words/2-grams to vectors, and then classify using a linear fully connected layer.

```
[0]: import torch.nn as nn
      import torch.nn.functional as F

      class TextSentiment(nn.Module):
          def __init__(self, vocab_size, embed_dim, num_class):
              super().__init__()
              self.embedding = nn.EmbeddingBag(vocab_size, embed_dim, sparse=True)
              self.fc = nn.Linear(embed_dim, num_class)
              self.init_weights()

          def init_weights(self):
              initrange = 0.5
              self.embedding.weight.data.uniform_(-initrange, initrange)
              self.fc.weight.data.uniform_(-initrange, initrange)
              self.fc.bias.data.zero_()

          def forward(self, text, offsets):
```

```

        embedded = self.embedding(text,offsets)
        return self.fc(embedded)

```

OK, now let's write a simple data loader and routine to train/test the model.

```

[0]: from torch.utils.data import DataLoader
embed_dim = 32
model = TextSentiment(vocab_size,embed_dim,num_class).to(device)

def train_func(sub_train):
    train_loss = 0
    train_acc = 0
    data =
    ↪DataLoader(sub_train,batch_size=BATCH_SIZE,shuffle=True,collate_fn=generate_batch)
    for i, (text,offsets,labels) in enumerate(data):
        optimizer.zero_grad()
        text,offsets,labels = text.to(device),offsets.to(device),labels.to(device)
        output = model(text,offsets)
        loss = criterion(output,labels)
        train_loss += loss.item()
        loss.backward()
        optimizer.step()
        train_acc += (output.argmax(1)==labels).sum().item()

    scheduler.step() #adjust learning rate

    return train_loss/len(sub_train), train_acc/len(sub_train)

def test_func(sub_test):
    test_loss = 0
    test_acc = 0
    data = DataLoader(sub_test,batch_size=BATCH_SIZE,collate_fn=generate_batch)
    for i, (text,offsets,labels) in enumerate(data):
        text,offsets,labels = text.to(device),offsets.to(device),labels.to(device)
        with torch.no_grad():
            output = model(text,offsets)
            loss = criterion(output,labels)
            test_loss += loss.item()
            test_acc += (output.argmax(1)==labels).sum().item()

    return test_loss/len(sub_test), test_acc/len(sub_test)

```

```

[0]: num_epochs = 5
criterion = torch.nn.CrossEntropyLoss().to(device)
optimizer = torch.optim.SGD(model.parameters(),lr=4)
scheduler = torch.optim.lr_scheduler.StepLR(optimizer,1,gamma=0.9)

```

```

for epoch in range(num_epochs):
    %time train_loss, train_acc = train_func(train_dataset)

    print('Epoch = %d' %(epoch+1))
    print(f'\tLoss: {train_loss:.4f}(train)\t|\tAcc: {train_acc * 100:.
→1f}% (train)')

```

```

CPU times: user 7.8 s, sys: 482 ms, total: 8.28 s
Wall time: 8.36 s
Epoch = 1
    Loss: 0.0259(train)      |      Acc: 84.8%(train)
CPU times: user 7.77 s, sys: 456 ms, total: 8.23 s
Wall time: 8.31 s
Epoch = 2
    Loss: 0.0118(train)      |      Acc: 93.7%(train)
CPU times: user 7.72 s, sys: 464 ms, total: 8.19 s
Wall time: 8.24 s
Epoch = 3
    Loss: 0.0070(train)      |      Acc: 96.3%(train)
CPU times: user 7.74 s, sys: 467 ms, total: 8.21 s
Wall time: 8.27 s
Epoch = 4
    Loss: 0.0040(train)      |      Acc: 98.0%(train)
CPU times: user 7.93 s, sys: 426 ms, total: 8.36 s
Wall time: 8.48 s
Epoch = 5
    Loss: 0.0023(train)      |      Acc: 99.0%(train)

```

```

[0]: %time test_loss, test_acc = test_func(test_dataset)
    print(f'\tLoss: {test_loss:.4f}(test)\t|\tAcc: {test_acc * 100:.1f}% (test)')

```

```

CPU times: user 201 ms, sys: 4.6 ms, total: 206 ms
Wall time: 207 ms
    Loss: 0.0238(test)      |      Acc: 90.3%(test)

```

Cool! Let us now see if we can test it on a real world example.

```

[0]: from torchtext.data.utils import ngrams_iterator
    from torchtext.data.utils import get_tokenizer

    def predict(text,model,vocab,ngrams):
        tk = get_tokenizer('basic_english')
        with torch.no_grad():
            text = torch.tensor([vocab[token] for token in_
→ngrams_iterator(tk(text),ngrams)])
            output = model(text,torch.tensor([0]))
            return output.argmax(1).item()

```

```
ex_string = """The world's airlines, no longer operating a globe-spanning_
↳choreography of flights,
    are consumed with new work: navigating government bailout offers, negotiating_
↳with unions,
    finding places to park idle planes and scrounging for business like flying_
↳cargo
    and repatriating marooned travelers."""

vocab = train_dataset.get_vocab()
model = model.to('cpu')
print('%s' %ag_news_label[predict(ex_string,model,vocab,2)])
```

Business