

# demo12

April 21, 2020

In this lab, we will learn how to implement k-means and EM, and visualize some essential differences. Let us start with some synthetic data.

```
[1]: %matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
/usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19:
FutureWarning: pandas.util.testing is deprecated. Use the functions in the
public API at pandas.testing instead.
import pandas.util.testing as tm
```

```
[0]: from sklearn.datasets import make_blobs
X, y_true = make_blobs(n_samples=400, centers=4, cluster_std=2, random_state=10)
```

```
[3]: plt.scatter(X[:, 0], X[:, 1], c=y_true, s=40, cmap='viridis');
ax = plt.gca()
ax.axis('equal')
```

```
[3]: (-11.02535175761933,
      11.24746483124909,
      -16.848150742263016,
      11.839765378970164)
```




OK, let's now run KMeans.

```
[4]: from sklearn.cluster import KMeans
kmeans = KMeans(4)
labels = kmeans.fit(X).predict(X)
plt.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap='viridis');
ax = plt.gca()
ax.axis('equal')
plt.show();
```




The first issue in k-means is that it does *hard* clustering: points at the borders between clusters get somewhat arbitrarily assigned, and there is no way to declare low cluster confidence. The second issue is that clusters are necessarily circular/spherical.

```
[5]: rng = np.random.RandomState(13)
X_stretched = np.dot(X, rng.randn(2, 2))
plt.scatter(X_stretched[:, 0], X_stretched[:, 1], c=y_true, s=40,
            cmap='viridis');
ax = plt.gca()
ax.axis('equal')
plt.show();
```




```
[6]: kmeans = KMeans(n_clusters=4, random_state=0)
      labels = kmeans.fit(X_stretched).predict(X_stretched)
      plt.scatter(X_stretched[:, 0], X_stretched[:, 1], c=labels, s=40,
                  cmap='viridis');
      ax = plt.gca()
      ax.axis('equal')
      plt.show();
```



Interesting. Because output cluster labels are forced to be circular, kmeans is picking up data points for different “true clusters”. This can be fixed via EM using a GMM model.

```
[7]: from sklearn.mixture import GaussianMixture as GMM
gmm = GMM(n_components=4)
labels = gmm.fit(X).predict(X)
plt.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap='viridis');
ax = plt.gca()
ax.axis('equal')
plt.show();
```




Slightly better clustering results, especially at the border. EM does “soft” clustering, so we can actually examine each point’s posterior probability.

```
[8]: probs = gmm.predict_proba(X)
      print(probs[:10].round(3))
```

```
[[0.  0.  0.  1.  ]
 [0.  1.  0.  0.  ]
 [0.  0.  0.  1.  ]
 [0.  0.  1.  0.  ]
 [0.  0.  1.  0.  ]
 [1.  0.  0.  0.  ]
 [1.  0.  0.  0.  ]
 [0.  0.996 0.  0.004]
 [0.  0.974 0.  0.026]
 [0.  0.999 0.  0.001]]
```

Let’s now test it on the stretched dataset.

```
[9]: labels = gmm.fit(X_stretched).predict(X_stretched)
      plt.scatter(X_stretched[:, 0], X_stretched[:, 1], c=labels, s=40,
                  cmap='viridis');
      ax = plt.gca()
      ax.axis('equal')
      plt.show();
```



Nice! The clusters have been currently identified.

(This part is optional). We can actually use GMM to fit lots of tiny clusters – this helps us to model non-spherical/elliptical datasets too.

```
[0]: from matplotlib.patches import Ellipse

def draw_ellipse(position, covariance, ax=None, **kwargs):
    """Draw an ellipse with a given position and covariance"""
    ax = ax or plt.gca()

    # Convert covariance to principal axes
    if covariance.shape == (2, 2):
        U, s, Vt = np.linalg.svd(covariance)
        angle = np.degrees(np.arctan2(U[1, 0], U[0, 0]))
        width, height = 2 * np.sqrt(s)
    else:
        angle = 0
        width, height = 2 * np.sqrt(covariance)

    # Draw the Ellipse
    for nsig in range(1, 4):
        ax.add_patch(Ellipse(position, nsig * width, nsig * height,
                              angle, **kwargs))

def plot_gmm(gmm, X, label=True, ax=None):
```

```

ax = ax or plt.gca()
labels = gmm.fit(X).predict(X)
if label:
    ax.scatter(X[:, 0], X[:, 1], c=labels, s=40, cmap='viridis', zorder=2)
else:
    ax.scatter(X[:, 0], X[:, 1], s=40, zorder=2)
ax.axis('equal')


w_factor = 0.2 / gmm.weights_.max()
for pos, covar, w in zip(gmm.means_, gmm.covariances_, gmm.weights_):
    draw_ellipse(pos, covar, alpha=w * w_factor)

```

```

[11]: from sklearn.datasets import make_moons
Xmoon, ymoon = make_moons(200, noise=.05, random_state=0)
plt.scatter(Xmoon[:, 0], Xmoon[:, 1]);

```




```

[12]: gmm16 = GMM(n_components=16, covariance_type='full', random_state=0)
plot_gmm(gmm16, Xmoon, label=False)

```





OK, let's do a real(istic) data experiment. Let's now use k-means to *automatically learn* MNIST digit classes without training labels.

```
[13]: from sklearn.datasets import load_digits
      digits = load_digits()
      digits.data.shape
```

```
[13]: (1797, 64)
```

```
[14]: kmeans = KMeans(n_clusters=10, random_state=0)
      clusters = kmeans.fit_predict(digits.data)
      kmeans.cluster_centers_.shape
```

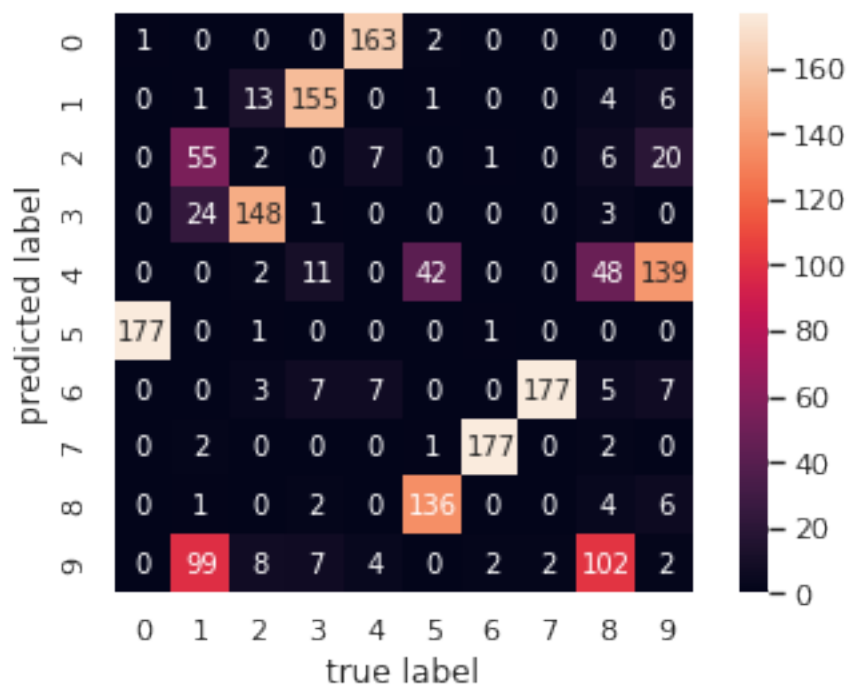
```
[14]: (10, 64)
```

```
[15]: fig, ax = plt.subplots(2, 5, figsize=(8, 3))
      centers = kmeans.cluster_centers_.reshape(10, 8, 8)
      for axi, center in zip(ax.flat, centers):
          axi.set(xticks=[], yticks=[])
          axi.imshow(center, interpolation='nearest', cmap=plt.cm.binary)
```



Quite remarkably, the cluster centers recover the digits (without ever telling the algorithm what they are). The digit “8” is a bit blurred out though.

```
[16]: from sklearn.metrics import confusion_matrix
mat = confusion_matrix(digits.target, clusters)
sns.heatmap(mat.T, square=True, annot=True, fmt='d',
             xticklabels=digits.target_names,
             yticklabels=digits.target_names)
plt.xlabel('true label')
plt.ylabel('predicted label');
```



Fairly accurate! We can compute accuracy by permuting the rows of the above matrix to match the max values of each row and column.

```
[17]: from scipy.stats import mode
corrected_labels = np.zeros_like(clusters)
for i in range(10):
    mask = (clusters == i)
    corrected_labels[mask] = mode(digits.target[mask])[0]

from sklearn.metrics import accuracy_score
accuracy_score(digits.target, corrected_labels)
```

```
[17]: 0.7952142459654981
```

Great! ~80% accuracy with no labels whatsoever!

Actually – if we do a bit more pre-processing, we can get even better results. There is a nonlinear dimensionality algorithm called (tSNE) which we don’t discuss – but which is very good at retaining “cluster” information. If we fit kmeans to a tSNE-reduced version of the dataset, we can improve performance all the way to 93% (which is even better than logistic regression/shallow neural nets) on MNIST. This example illustrates the power of unsupervised learning when used carefully.

```
[18]: from sklearn.manifold import TSNE

tsne = TSNE(n_components=2, init='random', random_state=0)
digits_proj = tsne.fit_transform(digits.data)

kmeans = KMeans(n_clusters=10, random_state=0)
clusters = kmeans.fit_predict(digits_proj)

labels = np.zeros_like(clusters)
for i in range(10):
    mask = (clusters == i)
    labels[mask] = mode(digits.target[mask])[0]

accuracy_score(digits.target, labels)
```

```
[18]: 0.9371174179187535
```