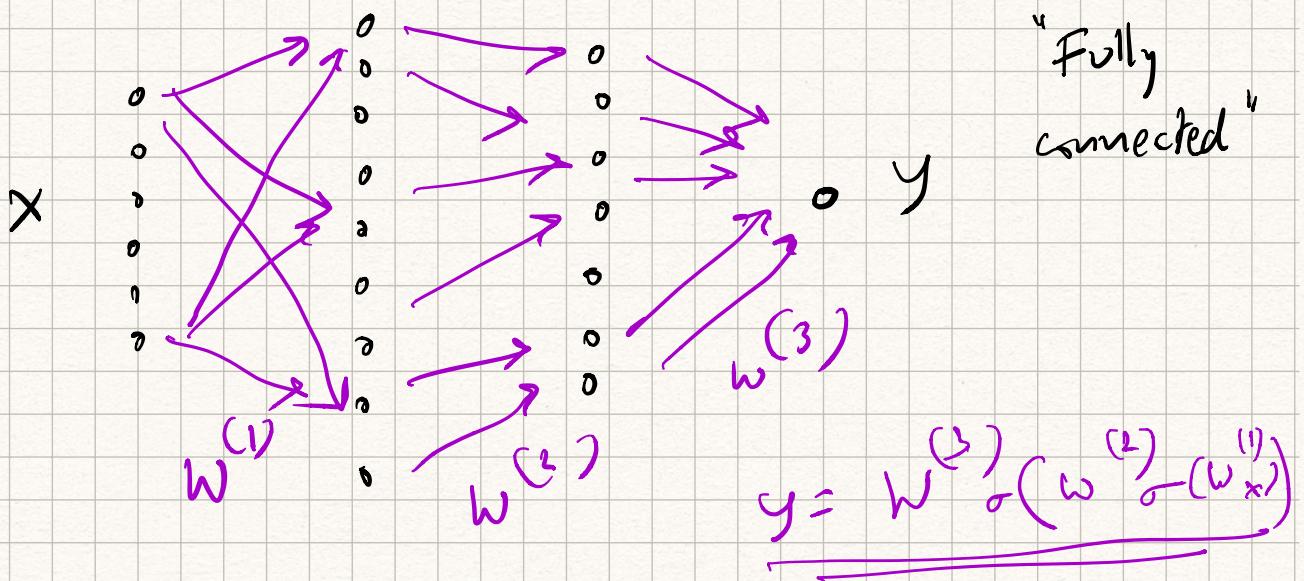


ECE6143 - GIY
Intro to ML

D Recap: Neural networks

D Convolutional Neural Nets (CNNs)

D Recurrent Neural Nets (RNNs)



Problems.

* Computational cost.

e.g. Image processing

Input : RGB images of 400×400

1 layer of neurons : 1000 neurons

Output : 10 classes.

edges (weights)

$$\approx 3 \times 400 \times 400 \times 1000 + 1000 \times 10$$

\approx 500 million weights.

Generally : $\Theta(L \cdot d^2)$

- Sample complexity.
 - Requires a very large training dataset.
- Loss of context
 - all input is flattened into vectors.
 - e.g. images lose spatial information
 - e.g. NLP / speech have long-range temporal dependencies.
- Confounding features.
 - Occlusion
 - Background clutter
 - Illumination
 - Other perturbations

Convolutional Neural Net. (CNNs) // image recognition

- Locality : the model should look at local features.

local regions of your input image.

- Shift-invariance : The model should make the ^{same} prediction to the same object no matter where it appears.

Convolutional layer

Convolution

$x[t]$ $w[t]$
input ↑ filter / kernel

$$\begin{aligned}y(t) &= (x * w)(t) \\&= \sum_{\tau} x[t-\tau] w(\tau)\end{aligned}$$

"Translate-and-scale".

No need to prove:

$$y(t) = \sum_{\tau} x(\tau) w[t-\tau].$$

"Flip-and-dot".

- Linear, associative, commutative

Shift-invariance //

2D-images

$$x(s, t)$$

$$w(s, t).$$

$$(x * w)(s, t) = \sum_{\sigma} \sum_{\tau} x[s - \sigma, t - \tau] w[\sigma, \tau]$$



Convolution layers

Same as convolution, except no flipping.

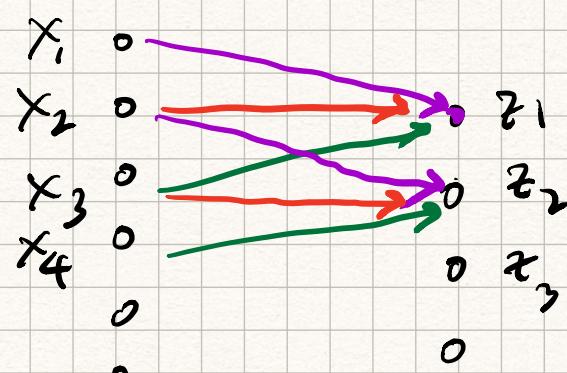
$$z(s, t) = (x * w)(s, t) = \sum_{\tau=-\Delta}^{\Delta} \sum_{\sigma=-\Delta}^{\Delta} x[s + \sigma, t + \tau] w[\sigma, \tau]$$

$$x[s + \sigma, t + \tau] \quad w[\sigma, \tau].$$

example $\Delta = 2$ \rightarrow restrict to $5 \times 5 = 25$ indices of input.

$$h(s, t) = \phi(z(s, t))$$

$$h(s, t) = \phi((x * w)(s, t))$$



Consequences:



1) Sparsely connected network.

2) Weight sharing.

In practice, we have D input channels.
 F output channels.

Input x_1, x_2, \dots, x_D .

Output z_1, z_2, \dots, z_F

$$\boxed{z_i = \sum_j x_j * w_{ij} \quad h_i = \phi(z_i)} \quad \text{↓}$$

Computational benefit:

$$(2\Delta + 1)^2 \cdot D \cdot F$$

$$\approx O(\Delta^2 \cdot D \cdot F).$$

Eg.

$$(2\Delta + 1)^2 = 25 \cdot 3 \cdot 10 \\ \approx 1000$$

- LeNet - 5
- Alex Net (2012)
- VGG Net (2014)
- Google Net (2014)
- Residual networks

Training : Back propagation.

Recurrent Neural Networks (Time series/
NLP-analysis).

NLP

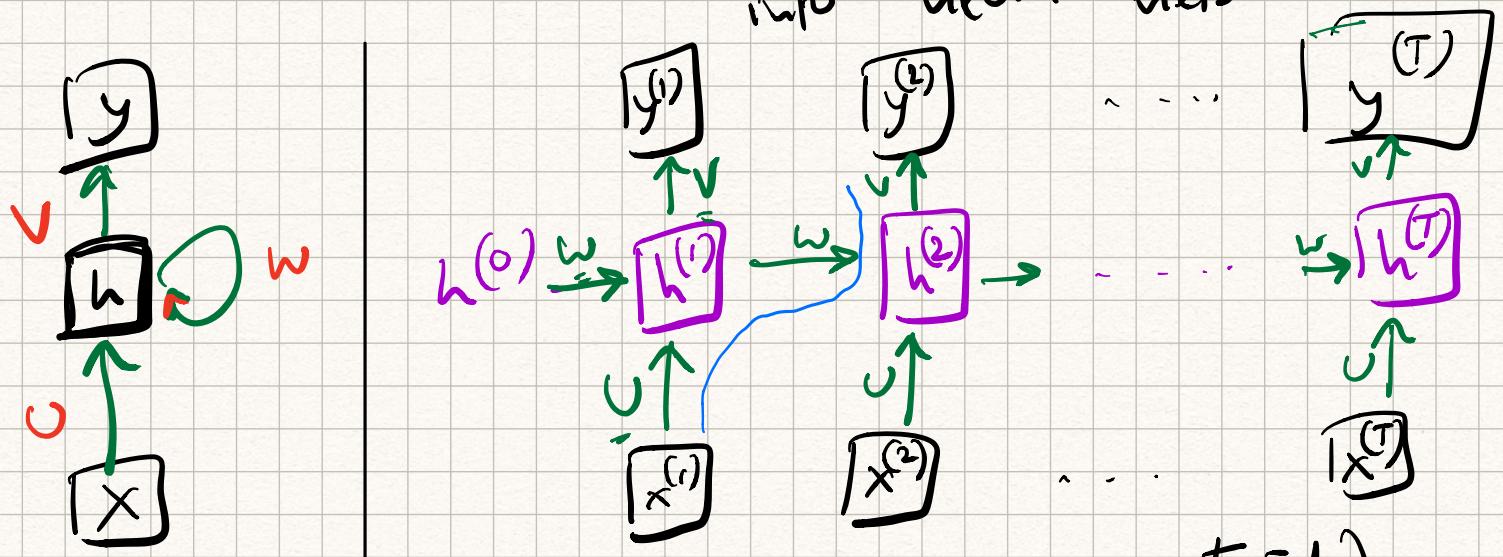
- Document retrieval
- Speech to text
- Language translation
- Sentiment prediction

- Short range dependencies.
- long range dependencies.

Classically: Markov models.

$$\begin{aligned}
 & P(\{w_1, w_2, \dots, w_d\}) \\
 & = p(w_1) p(w_2 | w_1) \underset{=}{=} p(w_3 | w_2) \dots \\
 & \qquad \qquad \qquad p(w_d | w_{d-1}) \\
 & \frac{}{p(w_3 | (w_1, w_2))} \\
 \end{aligned}$$

Recurrent neural nets: introduce feedback into neural nets.



$$\begin{aligned}
 \underline{h}^t &= \phi(Ux^t + Wh^{t-1}) // \\
 \underline{y}^t &= Vh^t
 \end{aligned}$$

. Observe: U, V, W don't change over time.
∴ Weight sharing.

Example:

- Basic RNN (h is single layer network).

- GRU (Gated recurrent units)
- LSTM (Long-short term memory).
- Attention networks.

